

Predicting Audio Features with Last.fm Tags

Jaime Ramírez Castillo^{1*} and M. Julia Flores^{1†}

¹Departamento de Sistemas Informáticos, Universidad de Castilla-La Mancha, Campus universitario s/n, Albacete, 02071, Spain.

*Corresponding author(s). E-mail(s):

Jaime.Ramirez@alu.uclm.es;

Contributing authors: Julia.Flores@uclm.es;

[†]These authors contributed equally to this work.

Abstract

In this paper, we discuss a number of experiments to analyze the suitability of music label representations to predict certain audio features, such as danceability, loudness, or acoustiness ...

Keywords: Music information retrieval, Artificial intelligence

1 Introduction

Music information retrieval (MIR) is an interdisciplinary research field that encompasses the extraction, processing, and knowledge discovery of information contained in music. MIR research intersects with other fields, such as computer science, signal processing, musicology, and sociology. The field covers applications in recommendation systems, music classification, music source separation, and music generation, among others [1].

MIR applications often attempt to extract information from the music audio signal, but analyzing associated metadata is also a common practice. Audio signals are typically preprocessed and transformed into intermediate formats, such as frequency-based signal representations or hand-crafted audio features. Associated metadata, such as editorial information, lyrics, or user-generated content is usually in text format. However, metadata can be available

as images of videos, for example, when analyzing album artwork, or music videos.

Depending on the specific MIR application, researchers or practitioners expect different output values. For example, an application that extracts audio features might return values for the tempo, the key, or the sample rate. In more complex applications, where, for example, machine learning is used, applications might return the estimated emotion that a track produces on a listener, or the predicted music genre of a particular track.

Among potentially useful input and output values, research has proved Spotify audio features and Last.fm tags to be significant values to characterize music. Spotify audio features capture high-level information about the music signal. Examples of these values are energy, danceability, or valence, among others. Last.fm tags are text labels that users associate to particular tracks, artists or albums in the Last.fm website.

Both types of values have been used as input values mostly for classification tasks, such as music genre recognition, where, given Spotify features or Last.fm tags, the model estimates the music genres of a particular track. However, these values have not been used as target values by previous research, to the best of our knowledge. In particular, the focus of this article is on predicting Spotify audio features, given a set of Last.fm tags.

By predicting Spotify audio features, we explore the relationship between subjective perception and concrete musical features. This approach might help to identify patterns and hidden correlations between how music is perceived, consumed and discovered.

Additionally, predicted audio features could be used in recommendation systems to provide users with explainable recommendations. Music recommendations are not always easy to interpret from the perspective of the listener. Users often get recommendations without meaningful explanations or justifications. Therefore, by using predicted Spotify features as basis for recommendation, we can explain users why the algorithm suggests a particular track. This process could be part of an explainable recommendation pipeline, where users enter a set of tags, and they receive the predicted audio features, the closest tracks to those features, and the distance values between each track and the predicted features.

In the remainder of the article, we explain the data gathering and preparation process, as well as the data input formats and various models. We will explore various models for the same track and provide insights on how accurately the prediction can be, by using only Last.fm tags.

2 Related Research

In recent years, researchers have studied the use of Last.fm tags in classification and regression tasks. Several studies have used Last.fm to predict music sentiment or mood.

In the last decade, Last.fm tags have been a popular source of metadata for MIR tasks. Last.fm tags can contain subjective information the genre, mood, and style of music, and might be use to characterize certain features of a music piece.

Last.fm tags can be useful when the audio signal available, for example, due to copyright limitations.

A number of studies have explored the use of Last.fm tags in MIR, and have shown promising results in predicting various audio features.

Researchers have used Last.fm tags. For example, Laurier et al. analyzed how Last.fm tags categorize mood. In their study, they created a semantic mood space based on Last.fm tags [2].

For example, Çano and Morisio discuss the process they follow to create a dataset of music lyrics annotated with Last.fm. In the creation process, they conclude that Last.fm tags are mostly related to music genre and positive moods [3]. In a similar direction, Bodó and Szilágyi generated a dataset for lyrics genre classification by combining the Last.fm with MusicBrainz data [4]. MusicBrainz ¹ is an online database of music editorial metadata. ¹

The Last.fm data has been the most widely used Last.fm dataset² in research. This dataset is a complementary dataset of the Million Song Dataset (MSD) [5].

Spotify, one of the leaders in the music streaming industry...

Additionally, the Spotify audio features have been used in multiple studies. Historically, these features were also called *Echo Nest audio features*. Spotify acquired Echo Nest in 2014 and made the audio features available via the Spotify API ³. Echonest was an online platform that was later acquired by Spotify.

Wang and Horvát use audio features to study differences between male and female artists [6]

In general, these studies confirm the possibility of extracting knowledge from Last.fm tags. To the best of our knowledge, no studies have addressed the problem of audio features regression, based solely on Last.fm tags.

Jamdar et al. used EchoNest audio features, combined with lyrics data to classify songs into emotion tags. These classes were first defined based on a Last.fm tags emotion mapping [7].

Similarly, Non-negative Matrix Factorization was applied in combination with EchoNest audio features for song recommendations [8].

P4kxspotify is a publicly available dataset that combines music review texts with Spotify audio features. The dataset creators argue that, although the terms of service prohibits scraping, their work is ethical [9].

In general, Spotify audio features have been used as predictive variables. We, to the best of our knowledge, are unaware of students that uses these features as target variables.

¹<https://musicbrainz.org/>

²Last.fm dataset, the official song tags and song similarity collection for the Million Song Dataset, available at: <http://millionsongdataset.com/lastfm>.

³https://en.wikipedia.org/wiki/The_Echo_Nest

While Spotify provides a description of the high-level audio features, how they compute or estimate these values is not publicly available. Panda and Redinho explore the use of Spotify high-level features applied to Music Emotion Recognition (MER) [10]. In particular, they identify that the energy, valence, and acousticness values, provided by the Spotify API, are highly relevant for emotion classification. They also achieve better performance on MER models by using their own top-100 features, and they determine that, although these three Spotify features are relevant in terms of characterizing emotion, more features are needed for MER.

3 Generating a Dataset

Before conducting experiments to predict audio features from tags, we constructed a dataset, by gathering the data from the Last.fm and Spotify APIs.

3.1 A Single-user Dataset

This work is scoped within our single-user research area [11]. In this area, we explore the development of music recommender systems that characterize the music preferences and listening context only for a single user. By training our system with single-user data, we also explore the following question: Is it possible to train recommender systems, and in particular, user-centric systems, by using a single-user dataset?

The user data for this work has been extracted from the listening history of the corresponding author user in Last.fm ⁴.

Similar to other intelligent systems, recommender systems must be trained, by using user preference data, to produce suitable recommendations. For this study, we leverage the knowledge discovery potential of large historical listening logs, gathered from Last.fm.

Given the objective of our experiments, we created a dataset of Last.fm tags and Spotify audio features, indexed by track, by following these steps:

3.2 Last.fm Tags

Last.fm is an online music service for users to keep track of their music listening habits. Last.fm is also an online community where users tag artists, albums, and tracks, according to their own taste and perception.

Users apply these tags to categorize music from their own perspective, which means that tags do not fit into any structured ontology or data model. Tags can refer to any aspect that users consider as a valid descriptor, such as genre, emotion, or user listening context.

⁴<https://www.last.fm/user/jimmydj2000>

For nearly two decades, users have been contributing to Last.fm by tagging tracks, albums, and artists with text labels. Although many of these descriptions are single-worded (e.g. *rock*, *dance*, or *happy*), users can also use short sentences to define a song, such as *I like this track*, or *on the beach*.

3.2.1 Last.fm Downloaded Data

Last.fm uses the term *scrobble* to refer to a single track playback, at a particular moment. We queried the Last.fm API to download the user's scrobble logs, reported from 2007 to 2022. For each scrobble, we have gathered the following information:

- Playback timestamp
- Track name
- Artist name
- Track tags. If the track does not have any tags assigned, then artist tags have been used

For each tag-track mapping, Last.fm includes a *count* value, which indicates the popularity of the given tag for the track. Last.fm normalizes this value in the 0-100 range, so the most popular tag for a track can have a count value of 100. For example, if *piano* is the most popular tag for a track, then the track might be probably associated to the following tuple (*piano*, 100).

Users normally listen to their favorite tracks several times, so the amount of unique tracks listened is much smaller than the number of track plays. In this case, the amount of individual tracks listened is about 20,000, and the number of scrobblings is, approximately, 90,000.

3.3 Spotify Audio Features

After gathering the listening history and track tags from Last.fm, and identifying the unique tracks that represent the user music collection, we collected Spotify audio features for each one of those tracks.

The Spotify audio features are numerical values that represent high-level audio information computed from a specific track. These values characterize a track, musically speaking, by measuring relevant musical aspects. For example, a danceability value of 0.95 means that a particular song is highly suitable for dancing.

The features provided by the Spotify API are listed in Table 1. The reader can find further details about each feature in the Spotify API documentation⁵.

The Spotify API failed to provide audio features for a portion of the tracks. These tracks were filtered out from our experiments. After filtering tracks that miss Last.fm tags or Spotify audio features, our dataset contains 14,009 samples.

⁵<https://developer.spotify.com/documentation/web-api/reference/#/operations/get-audio-features>

6 *Predicting Audio Features with Last.fm Tags***Table 1** Spotify audio features. These features provide high-level musical information about a track.

Feature name	Description
acousticness	The track is acoustic. From 0 to 1
danceability	The track encourages (or is adequate for) dancing. From 0 to 1
duration_ms	Duration in milliseconds
energy	The track is perceived as energetic. From 0 to 1
instrumentalness	The track is instrumental. From 0 to 1
key	Key categories encoded as integers. From C (0) to 11
liveness	The audience is audible. From 0 to 1
loudness	In decibels. From -60 to 0
mode	Major (1) or minor (0)
speechiness	Does the track contain speeches? From 0 to 1
tempo	In beats per minute (BPM)
valence	How happy is the track (BPM). From 0 to 1

Considering that the data was gathered from a single user, we explored the data to verify that the distribution of the Spotify audio features was comparable to other Spotify datasets. In particular, we verified that the distribution of the features, described in Table 2 and Figure 1, was comparable to the distribution of the *Spotify Audio Features* Kaggle dataset ⁶.

Table 2 Audio features description

Feature	μ	σ
Danceability	0.599	0.193
Energy	0.631	0.233
Acousticness	0.221	0.302
Instrumentalness	0.514	0.382
Valence	0.435	0.279

3.4 Filtering Missing Values

Not every track of the approximately unique tracks 20,000 of the Last.fm user' listening history included Last.fm tags and Spotify audio features. Tracks without Last.fm tags or without Spotify audio features were filtered out. As a result, of the originally 20,000 tracks included in the listening history, our dataset final size resulted in 14,009 samples.

4 Experiments

We trained commonly used machine learning models to predict an audio feature, given the set of tags for a particular track.

⁶<https://www.kaggle.com/datasets/tomigelo/spotify-audio-features>




Fig. 1 Distribution of audio features. These densities are similar to the values shown by the *Spotify Audio Features* Kaggle dataset.

- Boosted tree regressor [12]
- Naive Bayes Regressor [13]
- Fine-tuned GPT-2 model

The preceding models require specific input formats. Therefore, we tested different input formats.

4.1 Last.fm Tags Input Format

Each individual sample in the dataset corresponds to a unique track, and contains the list of Last.fm tag-count tuples (e.g. [(electronic, 100), (dance, 45), ...]) and the values of Spotify audio features. Before experimenting with machine learning models, we prepared the data in a number of different formats, each one suitable for specific models.

4.1.1 Tabular

With this format, the Last.fm tags are represented as a table. Each tag is defined by a column and each cell contains the count value of a tag for a track. A cell is 0 if a tag is missing for a track.

The number of columns is limited to the top-K tags. Counting the total amount of Last.fm tags in the user collection resulted, initially, in more than five million tags. We quickly confirmed that building a tabular data set, in which every row contains millions of columns (Last.fm tags) was doable, but presented scalability problems. Therefore, we reduced the number of tags by picking a subset of the most relevant tags.

The reduction algorithm is simple: calculate a weighted sum of all the tags appearances and pick the top 1000. The sum is weighted because we use the count attribute. This attribute is present in every Last.fm track-tag association and provides a measure of the strength of a particular tag in a

8 *Predicting Audio Features with Last.fm Tags*

specific track. Note that this reduction is an initial approach, which, similar to other phases, can be extended or improved in the future. For this particular case, dimensionality reduction algorithms, such as PCA, are good candidates for forthcoming iterations of this work.

The input data passed is formatted in tabular format, as follows:

- Given that $Tags_k$ is the set of most k frequent Last.fm tags in the user listening history and, where each $tag \in Tags_k$.
- Given that $Audio$ is the set of Spotify audio features, where each $feat \in Audio$.
- For each *track*:
 - $X_{track,tag}$ is the strength of *tag* for *track*. This value is in the 0–100 range.
 - $y_{track,feature}$ is the value of the audio feature *y* for *track*.

An example of this data format is provided in table 3.

Table 3 Tabular data format for Last.fm tags in XGBoost and Bayesian regressors

Track	$X_{electronic}$	$X_{ambient}$	$X_{...}$	y_{energy}	$y_{valence}$	$y_{...}$
Massive Attack - Blue Lines	62	6	...	0.496	0.947	...
The Beta Band - Squares	40	3	...	0.446	0.507	...
...

When generating training data by track, the tabular formats present sparsity problems.

For tabular representations, we need to defined a fixed set of columns as tags. For most of tracks, most columns are 0.

The sparsity of a matrix is the number of zero-valued elements divided by the total number of elements (e.g., $m * n$ for an $m * n$ matrix) is called the sparsity of the matrix.

4.1.2 Tabular Tokens

Tags are converted to text tokens. Columns represent token positions, and cells contain the token at a particular position, for a track. To tokenize tags, we have used the GTP2 tokenizer. Because the tokenizer requires a string as input, we have converted the set of tags for each track into a string. To *stringify* the tags, we have concatenated tags with multiple strategies:

- By including tag popularity: ‘rock 2, pop 1’.
- By repeating tags based on popularity: ‘rock rock, pop’.
- By ordering by popularity: ‘rock, pop’.

In this particular case, the X values of the tabular input data are tokens. These tokens are obtained from passing the a string of concatenated Last.fm tags through a tokenizer. The formal definition of this data format is as follows:

- Given that X_l is the token vocabulary, where l is the maximum vocabulary length.
- Given that $Audio$ is the set of Spotify audio features, where each $feat \in Audio$.
- For each $track$:
 - $X_{track,n}$ is token found at position n , after tokenizing the tags string.
 - $y_{track,feature}$ is the value of the audio feature y for $track$.

An example of this data format is provided in table 4.

Table 4 Tabular data format for tokens in XGBoost and Bayesian regressors

Track	X_0	X_1	X_2	X_{\dots}	y_{energy}	$y_{valence}$	y_{\dots}
Massive Attack - Blue Lines	101	5099	6154	...	0.496	0.947	...
The Beta Band - Squares	101	4522	2600	...	0.446	0.507	...
...

4.2 Text

When using transformer models, the input data is a string. We must represent the Last.fm tags, which are initially in the $(tagname, tagpopularity)$ form, to a string.

After converting to a string, the formal definition of the input data is as follows:

- Given that X is tags represented as text.
- Given that $Audio$ is the set of Spotify audio features, where each $feat \in Audio$.
- For each $track$:
 - $X_{track,n}$ is set of tags for $track$, encoded as a single string.
 - $y_{track,feature}$ is the value of the audio feature y for $track$.

An example of this data format is provided in table 5.

Table 5 Text data format for tokens in XGBoost and Bayesian regressors

Track	X	y_{energy}	$y_{valence}$	y_{\dots}
Massive Attack - Blue Lines	"hip hop, chill, bristol, ..."	0.496	0.947	...
The Beta Band - Squares	"alternative rock, folk, ..."	0.446	0.507	...
...

4.3 Boosted Tree Regressor

We configured the boosted tree regressor model with the training parameters listed in table 6.

Table 6 Training parameters for XGBoost regressor

Parameter	Value
objective	reg:squarederror
base score	0.5
booster	gbtree
colsample bylevel	1
colsample bynode	1
colsample bytree	1
gamma ¹	0
learning rate	0.300000012
max delta step	0
max depth	6
min child weight	1
estimators	200
n jobs	12
num parallel tree	1
predictor	auto
random state	0
reg alpha	0
reg lambda	1
scale pos weight	1
subsample	2
tree method	auto

¹Minimum loss reduction required to make a further partition on a leaf node of the tree.

4.4 Naive Bayes Regressor

The Naive Bayes Regressor, and in particular, Bayesian Ridge, is the model used for regression in this case.

The training parameters are listed in table 7.

4.5 Fine-tuned Transformer

Transformers are a type of neural network architecture that has been widely used in natural language processing (NLP) tasks, such as text generation and question answering. They were first introduced in 2017 by paper [REF: "Attention Is All You Need"] and have become one of the most popular models due to their ability to handle long-range dependencies and process variable-length inputs.

One popular variant of transformers is the Generative Pre-trained Transformer 2 (GPT-2), which was introduced by OpenAI in 2019. GPT-2 is a

Table 7 Training parameters for XGBoost regressor

Parameter	Value
Maximum iterations	300
Tolerance ¹	1×10^{-3}
alpha 1	1×10^{-6}
alpha 2	1×10^{-6}
lambda 1	1×10^{-6}
lambda 2	1×10^{-6}

¹Tolerance for the stopping criteria.

language model that has been trained on a large corpus of text, and can generate coherent and fluent text that resembles human-written language. GPT-2 has achieved remarkable results in a wide range of natural language processing tasks, including text generation, machine translation, and question answering.

In this study, we have used GPT-2 as a regressor to predict Spotify audio features from Last.fm tags. The basic idea behind this approach is to feed a string of concatenated Last.fm tags as input to the GPT-2 model, and then use the model's output as the predicted value of a specific Spotify audio feature. By training the GPT-2 model on a large dataset of Last.fm tags and corresponding Spotify audio features, we aim to learn the complex relationships between these two types of data and to use this knowledge to evaluate the prediction accuracy of Spotify audio features based on Last.fm tags.

To train the GPT-2 model as a regressor, we used a supervised learning approach. Specifically, we collected a large dataset of Last.fm tags and corresponding Spotify audio features, and used this dataset to train the GPT-2 model to predict Spotify audio features from Last.fm tags. We used a mean squared error loss function to optimize the model's performance during training, and we also employed techniques such as early stopping and learning rate scheduling to prevent overfitting and improve generalization performance.

Under development...

4.6 Experiments Execution and Results

Table 8 summaries the experiment results. The table provides RMSE values for each experiment.

4.6.1 Results for Tabular Data Models

5 Conclusions

In general, we believe that this novel approach that has the potential to benefit both listeners and researchers. By combining subjective user-generated tags with objective audio features, we can gain new insights into the complex relationship between perception and audio signal in music.

Table 8 Experiment results. Cells values correspond to the RMSE value.

M	Input format	Danceab...	Acoustic...	Energy	Valence	Instrumen...
Base		0.276	0.438	0.329	0.395	0.541
Bayes	100 tags ¹	0.159	0.261	0.197	0.243	0.307
Tree	100 tags	0.154	0.257	0.188	0.240	0.302
Bayes	1000 tags	0.153	0.253	0.190	0.237	0.299
Tree	1000 tags	0.149	0.249	0.184	0.236	0.291
Bayes	10000 tags	0.152	0.251	0.189	0.236	0.297
Tree	10000 tags	0.148	0.250	0.181	0.236	0.290
Bayes	100 tokens D ²	0.307	0.307	0.238	0.281	0.383
Tree	100 tokens D	0.274	0.274	0.212	0.256	0.330
Bayes	1000 tokens D	0.201	0.315	0.249	0.248	0.399
Tree	1000 tokens D	0.173	0.278	0.215	0.220	0.339
Bayes	10000 tokens D	0.359	0.507	0.394	0.479	0.613
Tree	10000 tokens D	0.172	0.276	0.217	0.266	0.342
Bayes	100 tokens O ³	0.191	0.305	0.237	0.282	0.376
Tree	100 tokens O	0.179	0.294	0.225	0.271	0.350
Bayes	1000 tokens O	0.237	0.343	0.276	0.339	0.428
Tree	1000 tokens O	0.182	0.294	0.224	0.270	0.353
Bayes	10000 tokens O	0.237	0.343	0.276	0.339	0.428
Tree	10000 tokens O	0.182	0.294	0.225	0.267	0.353
Bayes	100 tokens TC ⁴	0.191	0.304	0.236	0.281	0.380
Tree	100 tokens TC	0.172	0.280	0.211	0.262	0.342
Bayes	1000 tokens TC	0.202	0.320	0.247	0.248	0.404
Tree	1000 tokens TC	0.174	0.282	0.215	0.220	0.344
Bayes	10000 tokens TC	0.234	0.341	0.274	0.321	0.430
Tree	10000 tokens TC	0.175	0.281	0.214	0.270	0.345
.....
GPT	Duplicated ²	0.157
GPT	Ordered ³	0.149
GPT	Tags,Counts ⁴	0.145

¹Tags in tabular format. Given $(rock, 3)$, the cell in the *rock* column contains 3.

²Duplicated tokens in tabular format. Tags $(rock, 3)$, $(pop, 2)$, are converted to the "**rock, rock, rock, pop, pop**" string, which a tokenizer converts to the list of input tokens (e.g. [101,1005,16588,1005,2531, ...]). These tokens are passed to the model in tabular format. Columns are $token_1, token_2, \dots, token_N$.

³Ordered tokens in tabular format. Tags $(rock, 3)$, $(pop, 2)$, are converted to the "**rock, pop**" string, which a tokenizer converts to the list of input tokens (e.g. [101,1005,16588,1005,2531, ...]). These tokens are passed to the model in tabular format. Columns are $token_1, token_2, \dots, token_N$.

⁴Tokens in tabular format from tags and counts. Tags $(rock, 3)$, $(pop, 2)$, are converted to the "**'rock' 3, 'pop' 2**" string, which a tokenizer converts to the list of input tokens (e.g. [101,1005,16588,1005,2531, ...]). These tokens are passed to the model in tabular format. Columns are $token_1, token_2, \dots, token_N$.

²String. Given tags $(rock, 3)$, $(pop, 2)$, input is formatted as "**rock, rock, rock, pop, pop**".

³String. Given tags $(rock, 3)$, $(pop, 2)$, input is formatted as "**rock, pop**".

⁴String. Given tags $(rock, 3)$, $(pop, 2)$, input is formatted as "**'rock' 3, 'pop' 2**".

Our approach also presents limitations. One limitation is the assumption of a strong relationship between Last.fm tags and Spotify features, which may




Fig. 2 RMSE mean and standard deviation by model and tags/tokens limit.

not be true in all cases. Future work could explore other sources of input values, possibly related to the user context, to improve the accuracy of the predictions.

TODO

6 Acknowledgments

TODO

References

- [1] Ramirez, J., Flores, M.J.: Machine learning for music genre: multi-faceted review and experimentation with audioset. *Journal of Intelligent Information Systems* **55**(3), 469–499 (2020)
- [2] Laurier, C., Sordo, M., Serra, J., Herrera, P.: Music mood representations from social tags. In: *ISMIR*, pp. 381–386 (2009)
- [3] Çano, E., Morisio, M., *et al.*: Music mood dataset creation based on last. fm tags. In: *International Conference on Artificial Intelligence and Applications*, Vienna, Austria, pp. 15–26 (2017)




Fig. 3 RMSE mean and standard deviation by model and audio feature.

- [4] Bodó, Z., Szilágyi, E.: Connecting the last. fm dataset to lyricwiki and musicbrainz. lyrics-based experiments in genre classification. *Acta Universitatis Sapientiae, Informatica* **10**(2), 158–182 (2018)
- [5] Bertin-Mahieux, T., Ellis, D.P.W., Whitman, B., Lamere, P.: The million song dataset. In: *Proceedings of the 12th International Conference on Music Information Retrieval (ISMIR)* (2011)
- [6] Wang, Y., Horvát, E.-Á.: Gender differences in the global music industry: Evidence from musicbrainz and the echo nest. In: *Proceedings of the International AAAI Conference on Web and Social Media*, vol. 13, pp. 517–526 (2019)
- [7] Jamdar, A., Abraham, J., Khanna, K., Dubey, R.: Emotion analysis of songs based on lyrical and audio features. *arXiv preprint arXiv:1506.05012* (2015)
- [8] Benzi, K., Kalofolias, V., Bresson, X., Vandergheynst, P.: Song recommendation with non-negative matrix factorization and graph total variation. In: *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 2439–2443 (2016). Ieee
- [9] Pinter, A.T., Paul, J.M., Smith, J., Brubaker, J.R.: P4kxspotify: A




Fig. 4 RMSE mean and standard deviation by model and input type (tag probabilities or tokens).

- dataset of pitchfork music reviews and spotify musical features. In: Proceedings of the International AAAI Conference on Web and Social Media, vol. 14, pp. 895–902 (2020)
- [10] Panda, R., Redinho, H., Gonçalves, C., Malheiro, R., Paiva, R.P.: How does the spotify api compare to the music emotion recognition state-of-the-art? In: 18th Sound and Music Computing Conference (SMC 2021), pp. 238–245 (2021)
- [11] Ramirez-Castillo, J., Flores, M.J., Nicholson, A.E.: User-centric music recommendations. In: 16th Bayesian Modelling Applications Workshop, Conference on Uncertainty in Artificial Intelligence, Eindhoven, The Netherlands) (2022)
- [12] xgboost: Xgboost. <https://xgboost.readthedocs.io/en/stable/tutorials/model.html>
- [13] Tipping, M.E.: Sparse bayesian learning and the relevance vector machine. J. Mach. Learn. Res. **1**, 211–244 (2001)




Fig. 5 RMSE mean and standard deviation by audio feature and input type (tag probabilities or tokens).