

Predicting Spotify Audio Features from Last.fm Tags

Jaime Ramírez Castillo^{1*} and M. Julia Flores^{1†}

¹Departamento de Sistemas Informáticos, Universidad de Castilla-La Mancha, Campus universitario s/n, Albacete, 02071, Spain.

*Corresponding author(s). E-mail(s):

Jaime.Ramirez@alu.uclm.es;

Contributing authors: Julia.Flores@uclm.es;

[†]These authors contributed equally to this work.

Abstract

In this paper, we discuss a number of experiments to analyze the suitability of music label representations to predict certain audio features, such as danceability, loudness, or acoustiness ...

Keywords: Music information retrieval, Artificial intelligence

1 Introduction

Music information retrieval (MIR) is an interdisciplinary research field that encompasses the extraction, processing, and knowledge discovery of information contained in music. MIR research intersects with other fields, such as computer science, signal processing, musicology, and sociology. The field covers applications in recommendation systems, music classification, music source separation, and music generation, among others [1].

MIR applications often attempt to extract information from the music audio signal, but analyzing associated metadata is also a common practice. Audio signals are typically preprocessed and transformed into intermediate formats, such as frequency-based signal representations or hand-crafted audio

features. Associated metadata, such as editorial information, lyrics, or user-generated content is usually in text format. However, metadata can be available as images of videos, for example, when analyzing album artwork, or music videos.

Depending on the specific MIR application, researchers or practitioners expect different output values. For example, an application that extracts audio features might return values for the tempo, the key, or the sample rate. In more complex applications, where, for example, machine learning is used, applications might return the estimated emotion that a track produces on a listener, or the predicted music genre of a particular track.

Among potentially useful input and output values, research has proved Spotify audio features and Last.fm tags to be significant values to characterize music. Spotify audio features capture high-level information about the music signal. Examples of these values are energy, danceability, or valence, among others. Last.fm tags are text labels that users associate to particular tracks, artists or albums in the Last.fm website.

Both types of values have been used as input data mostly for classification tasks, such as music genre recognition, where, given Spotify features or Last.fm tags, the model estimates the music genres of a particular track. Previous studies, however, have not used these values as target outputs, to the best of our knowledge. In particular, the focus of this article is on predicting Spotify audio features, given a set of Last.fm tags.

By predicting Spotify audio features, we explore the relationship between subjective perception and concrete musical features. This approach might help to identify patterns and hidden correlations between how music is perceived, consumed and discovered.

Additionally, predicted audio features could be used in recommendation systems to provide users with explainable recommendations. Music recommendations are typically difficult to interpret from the perspective of the listener. Users often get recommendations without meaningful explanations or justifications. Therefore, by using predicted Spotify features as basis for recommendation, we can explain users why the algorithm suggests a particular track. This process could be part of an explainable recommendation pipeline, where users enter a set of tags, and they receive the predicted audio features, the closest tracks to those features, and the distance values between each track and the predicted features.

In the remainder of the article, we explain the data gathering and preparation process, as well as the data input formats and various models. We will explore various models for the same track and provide insights on how accurately the prediction can be, by using only Last.fm tags.

2 Related Research

2.0.1 Last.fm Tags

In the last decade, researchers have studied the use of Last.fm tags in classification and regression tasks. Last.fm tags have been a popular source of metadata for MIR tasks, because they potentially contain subjective information related to the genre, mood, and style of music, and might be used to characterize certain features of a music piece. Several studies have used Last.fm to predict music sentiment or mood, and to predict audio features.

Additionally, Last.fm tags constitute a useful source of input knowledge when the audio signal is not available, for example, due to copyright limitations.

For example, Laurier et al. analyze how Last.fm tags categorize mood. In their study, they created a semantic mood space based on Last.fm tags [2].

Čano and Morisio discuss how they create a dataset of music lyrics annotated with Last.fm tags. In the creation process, they conclude that Last.fm tags are mostly related to music genre and positive moods [3].

In a similar direction, Bodó and Szilágyi generate a dataset for lyrics genre classification by combining the Last.fm with MusicBrainz data [4]. MusicBrainz is an online database of music editorial metadata¹.

So far, The *Last.fm dataset* has been the most widely used Last.fm dataset² in research. This dataset is a complementary dataset of the Million Song Dataset (MSD) [5].

In general, these studies confirm the possibility of extracting knowledge from Last.fm tags.

2.0.2 Spotify Audio Features

The Spotify audio features have been used in multiple studies. Historically, these features were called the *Echo Nest audio features*. The *Echo Nest* was an online music intelligence platform that provided users and clients with music analysis services. Among these services, the Echo Nest offered a database and an API to retrieve audio features for each of the tracks in the database³. Spotify acquired the Echo Nest in 2014. As a result, the Echo Nest API was shut down and Spotify made these audio features available via the Spotify API.

Nowadays, published research is available under the two terms. Whereas the most recent studies refer to these audio features as the Spotify audio features, earlier studies use the Echo Nest denomination.

Regardless of the term used, the features remain the same. These features are a set of high-level descriptors, such as energy and danceability, which are related to the audio but also to the listeners perception.

¹<https://musicbrainz.org/>

²Last.fm dataset, the official song tags and song similarity collection for the Million Song Dataset, available at: <http://millionsongdataset.com/lastfm>.

³https://en.wikipedia.org/wiki/The_Echo_Nest

While Spotify provides a description of the audio features, how they compute or estimate these values is not publicly available.

Wang and Horvát use audio features to study differences between male and female artists [6].

Jamdar et al. used EchoNest audio features, combined with lyrics data to classify songs into emotion tags. These classes were first defined based on a Last.fm tags emotion mapping [7].

Similarly, Non-negative Matrix Factorization was applied in combination with EchoNest audio features for song recommendations [8].

Panda and Redinho explore the use of Spotify high-level features applied to Music Emotion Recognition (MER) [9]. In particular, they identify that the energy, valence, and acousticness values, provided by the Spotify API, are highly relevant for emotion classification. They also achieve better performance on MER models by using their own top-100 features, and they determine that, although these three Spotify features are relevant in terms of characterizing emotion, more features are needed for MER.

Publicly available datasets of Spotify audio features can be found online, as a result of open-source and research communities collecting data from the Spotify API and publishing the results. It is unclear, however, whether these published datasets violate the Spotify API terms of service.

For example, *P4kxspotify* is a publicly available dataset that combines music review texts with Spotify audio features. The dataset creators argue that, although the terms of service prohibits scraping, their work is ethical [10].

Another example is the *Spotify Audio Features* Kaggle dataset⁴. This dataset contains more than 116,000 unique tracks, and includes audio features for each track.

In general, Spotify audio features have been used as predictive input variables. We, to the best of our knowledge, are unaware of studies that use these features as target variables, or studies that have addressed the problem of audio features regression, based solely on Last.fm tags.

3 Generating a Dataset

Before conducting experiments to predict audio features from tags, we constructed a dataset, by gathering the data from the Last.fm and Spotify APIs.

3.1 A Single-user Dataset

This work is scoped within our single-user research line [11]. In this area, we explore the development of music recommender systems that characterize the music preferences and listening context only for a single user. By training our system with single-user data, we also raise the following question: Is it possible to train recommender systems, and in particular, user-centric systems, by using a single-user dataset?

⁴<https://www.kaggle.com/datasets/tomigelo/spotify-audio-features>

The user data for this work has been extracted from the listening history of the corresponding author user in Last.fm⁵.

Similar to other intelligent systems, recommender systems must be trained, by using user preference data, to produce suitable recommendations. For this study, we leverage the knowledge discovery potential of large historical listening logs, gathered for the mentioned user from Last.fm.

Given the objective of our experiments, we created a dataset of Last.fm tags and Spotify audio features, indexed by track, by following these steps:

3.2 Last.fm Tags

Last.fm is an online music service for users to keep track of their music listening habits. Last.fm is also an online community where users tag artists, albums, and tracks, according to their own taste and perception.

Users apply these tags to categorize music from their own perspective, which means that tags do not fit into any structured ontology or data model. Tags can refer to any aspect that users consider as a valid descriptor, such as genre, emotion, or user listening context.

For nearly two decades, users have been contributing to Last.fm by tagging tracks, albums, and artists with text labels. Although many of these descriptions are single-worded (e.g. *rock*, *dance*, or *happy*), users can also use short sentences to define a song, such as *I like this track*, or *on the beach*.

3.2.1 Last.fm Downloaded Data

Last.fm uses the term *scrobble* to refer to the action of playing a track at a specific moment in time. Initially, Last.fm monitored user listening activity with a desktop application called *Scrobbler*. Users used to install this application on their computers to monitor their activity on players such as Winamp or iTunes. With the advent of music streaming services, the possibilities for users to scrobble their music habits expanded. Integrations were developed to integrate the scrobbler into popular platforms, such as Spotify, YouTube, or SoundCloud. Mobile versions of the Scrobbler were developed for Android and iOS devices, and also open source initiatives flourished⁶.

For us, the first step to construct the dataset was to download the user listening activity. We queried the Last.fm API to download the user's scrobbling logs, reported from 2007 to 2022. For each scrobble, we have gathered the following information:

- Playback timestamp.
- Track name.
- Artist name.
- Track tags. If the track does not have any tags assigned, then artist tags were used.

⁵<https://www.last.fm/user/jimmydj2000>

⁶<https://github.com/elamperti/OpenWebScrobbler>

Last.fm maps each track (and artist) to a list of community-contributed tags. For each track-tag mapping, Last.fm includes a *count* value, which indicates the popularity of the given tag for the track. Last.fm normalizes this value in the 0-100 range, so the most popular tag for a track can have a count value of 100. For example, if *jazz club* is the most popular tag for a track, then the track might be probably associated to the following tuple (*jazz club*, 100).

Users typically listen to their favorite tracks several times, so the amount of unique tracks played is much smaller than the number of track plays. In this case, the amount of individual tracks listened is about 20,000 and the number of scrobbles is, approximately, 90,000. Therefore, the user has listened to each song, approximately, 4.5 times on average.

3.3 Spotify Audio Features

After gathering the listening history and track tags from Last.fm, and identifying the unique tracks that represent the user music collection, we collected Spotify audio features for each one of those tracks.

The Spotify audio features are numerical values that represent high-level audio information computed from a specific track. These values characterize a track, musically speaking, by measuring relevant musical aspects. For example, a danceability value of 0.95 means that a particular song is highly suitable for dancing.

The features provided by the Spotify API are listed in Table 1. The reader can find further details about each feature in the Spotify API documentation⁵.

Table 1 Spotify audio features. These features provide high-level musical information about a track.

Feature name	Description
acousticness	The track is acoustic. From 0 to 1
danceability	The track encourages (or is adequate for) dancing. From 0 to 1
duration_ms	Duration in milliseconds
energy	The track is perceived as energetic. From 0 to 1
instrumentalness	The track is instrumental. From 0 to 1
key	Key categories encoded as integers. From C (0) to 11
liveness	The audience is audible. From 0 to 1
loudness	In decibels. From -60 to 0
mode	Major (1) or minor (0)
speechiness	Does the track contain speeches? From 0 to 1
tempo	In beats per minute (BPM)
valence	How happy is the track (BPM). From 0 to 1

⁵<https://developer.spotify.com/documentation/web-api/reference/#/operations/get-audio-features>

3.4 Filtering Missing Values

After gathering the data, we identified that the Spotify API had failed to provide audio features for a portion of the tracks. Similarly, Last.fm also returned no tags for another subset of the tracks. To prevent problems with missing values, we decided to filter out these tracks from the dataset, and therefore, from our experiments. After filtering tracks that were missing Last.fm tags or Spotify audio features, the dataset resulted in 14,009 samples. Compared to the original 20,000 unique tracks included in the listening history, approximately 6,000 songs were missing either Spotify or Last.fm data.

3.5 Dataset Comparison

Considering that the data was gathered from a single user, we explored the data to verify that the distribution of the Spotify audio features was comparable to larger, and possibly more balanced, Spotify datasets. In particular, we verified that the distribution of the features, described in Table 2 and Figure 1, was comparable to the distribution of the *Spotify Audio Features* Kaggle dataset. Our dataset presents similar μ and σ values to the Spotify Audio Features Kaggle dataset, as illustrated in Table 3 and Figure 2.

Table 2 Audio features description

Feature	μ	σ
Danceability	0.60	0.19
Energy	0.63	0.23
Acousticness	0.22	0.30
Instrumentalness	0.51	0.38
Valence	0.44	0.28

Table 3 Audio features description of *Spotify Audio Features* Kaggle dataset

Feature	μ	σ
Danceability	0.58	0.19
Energy	0.57	0.26
Acousticness	0.34	0.25
Instrumentalness	0.22	0.36
Valence	0.44	0.26

4 Experiments

We trained commonly used machine learning models to predict Spotify audio features, given the set of Last.fm tags for a particular track.

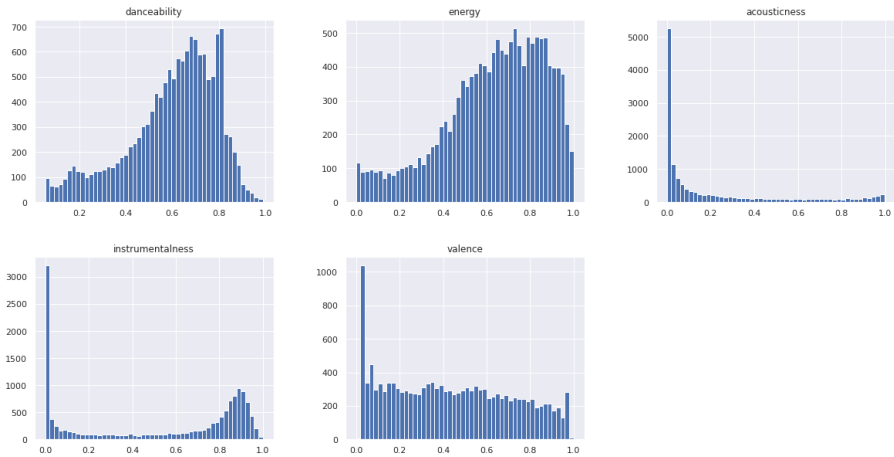
8 *Predicting Spotify Audio Features from Last.fm Tags*

Fig. 1 Distribution of audio features in our single-user dataset.




Fig. 2 Distribution of audio features in the *Spotify Audio Features* Kaggle dataset.

- Boosted tree regressor [12]
- Naive Bayes Regressor [13]
- Fine-tuned GPT-2 model

The preceding models require specific input formats. Therefore, we tested different input formats.

4.1 Last.fm Tags Input Format

Each individual sample in the dataset corresponds to a unique track, and contains the list of Last.fm tag-count tuples (e.g. [(electronic, 100), (dance, 45), ...]) and the values of Spotify audio features. Before experimenting with machine learning models, we prepared the data in a number of different formats, each one suitable for specific models.

4.1.1 Tabular

With this format, the Last.fm tags are represented as a table. Each tag is defined by a column and each cell contains the count value of a tag for a track. A cell is 0 if a tag is missing for a track.

Counting the total amount of Last.fm tags in the user collection resulted, initially, in more than five million tags. We quickly confirmed that building a tabular data set, in which every row contains millions of columns (Last.fm tags) was doable, but presented scalability problems. Therefore, we reduced the number of tags by picking a subset of the most relevant tags.

The reduction algorithm uses a basic data aggregation algorithm: group the data by tag, aggregate by summing the *count* values for each tag, order by the aggregated count, and finally select the top- K items of the ranking.

We generated three versions of the dataset, with different values of K : 100, 1,000 and 10,000.

We consider this reduction approach an initial approach. For this particular aspect, dimensionality reduction algorithms, such as PCA, are good candidates for future work.

After selecting the top- K tags, the input data passed was formatted in tabular format, as follows:

- Given that $Tags_K$ is the set of most K frequent Last.fm tags in the user listening history and, where each $tag \in Tags_K$.
- Given that $Audio$ is the set of Spotify audio features, where each $feat \in Audio$.
- For each $track$:
 - $X_{track,tag}$ is the strength of tag for $track$. This value is in the 0–100 range.
 - $y_{track,feature}$ is the value of the audio feature y for $track$.

An example of this data format is provided in table 4.

Table 4 Tabular data format for Last.fm tags in XGBoost and Bayesian regressors

Track	$X_{electronic}$	$X_{ambient}$	$X_{...}$	y_{energy}	$y_{valence}$	$y_{...}$
Massive Attack - Blue Lines	62	6	...	0.496	0.947	...
The Beta Band - Squares	40	3	...	0.446	0.507	...
...

When generating training data by track, the tabular formats present sparsity problems.

For tabular representations, we need to defined a fixed set of columns as tags. For most of tracks, most columns are 0.

The sparsity of a matrix is the number of zero-valued elements divided by the total number of elements (e.g., $m * n$ for an $m * n$ matrix) is called the sparsity of the matrix.

4.1.2 Tabular Tokens

Tags are converted to text tokens. Columns represent token positions, and cells contain the token at a particular position, for a track. To tokenize tags, we have used the GTP2 tokenizer. Because the tokenizer requires a string as input, we have converted the set of tags for each track into a string. To *stringify* the tags, we have concatenated tags with multiple strategies:

- Ordering by count: "rock, pop".
- Including tag count: "rock 2, pop 1".
- Duplicating each tag *count* times: "rock rock, pop".

In this particular case, the X values of the tabular input data are tokens. These tokens are obtained from passing the string of concatenated Last.fm tags through the GPT-2 tokenizer, as the following procedure explains:

- Given that X_L is the token vocabulary, where L is the maximum vocabulary length.
- Given that *Audio* is the set of Spotify audio features, where each $feat \in \text{Audio}$.
- For each *track*:
 - $X_{track,n}$ is token found at position n , after tokenizing the tags string.
 - $y_{track,feature}$ is the value of the audio feature y for *track*.

An example of this data format is provided in table 5.

Table 5 Tabular data format for tokens in XGBoost and Bayesian regressors

Track	X_0	X_1	X_2	X_{\dots}	y_{energy}	$y_{valence}$	y_{\dots}
Massive Attack - Blue Lines	101	5099	6154	...	0.496	0.947	...
The Beta Band - Squares	101	4522	2600	...	0.446	0.507	...
...

Similarly to the tags tabular format, we also defined fixed values for the number of columns: 10, 1,000, and 10,000.

The two tabular formats, with tags and tokens, were tested on the Naive Bayes and the Boosted tree regressors.

4.2 Text Strings

When using transformer models, the input data is a string. We must represent the Last.fm tags, which are defined as $(tag, count)$ tuples, as strings. To this end, we applied the same three transformations used in the tabular tokens formats, concatenate the tags by ordering by count, by including the tag count, and by duplicating the tags.

After converting to a string, the formal definition of the input data is as follows:

- Given that X is tags represented as text.
- Given that $Audio$ is the set of Spotify audio features, where each $feat \in Audio$.
- For each $track$:
 - $X_{track,n}$ is set of tags for $track$, encoded as a single string.
 - $y_{track,feature}$ is the value of the audio feature y for $track$.

An example of this data format is provided in table 6.

Table 6 Text data format for GPT-transformer. In this particular case, the tags have been concatenated by ordering by tag count

Track	X	y_{energy}	$y_{valence}$	$y...$
Massive Attack - Blue Lines	"hip hop, chill, bristol, ..."	0.496	0.947	...
The Beta Band - Squares	"alternative rock, folk, ..."	0.446	0.507	...
...

4.3 Boosted Tree Regressor

We configured the boosted tree regressor model with the training parameters listed in table 7.

4.4 Naive Bayes Regressor

The Naive Bayes Regressor, and in particular, Bayesian Ridge, is the model used for regression in this case.

The training parameters are listed in table 8.

4.5 Fine-tuned Transformer

Transformers are a type of neural network architecture that has been widely used in natural language processing (NLP) tasks, such as text generation and question answering. They were first introduced in 2017 by paper [REF: "Attention Is All You Need"] and have become one of the most popular models due to their ability to handle long-range dependencies and process variable-length inputs.

One popular variant of transformers is the Generative Pre-trained Transformer 2 (GPT-2), which was introduced by OpenAI in 2019. GPT-2 is a language model that has been trained on a large corpus of text, and can generate coherent and fluent text that resembles human-written language. GPT-2 has achieved remarkable results in a wide range of natural language processing tasks, including text generation, machine translation, and question answering.

In this study, we have used GPT-2 as a regressor to predict Spotify audio features from Last.fm tags. The basic idea behind this approach is to feed a string of concatenated Last.fm tags as input to the GPT-2 model, and then use

Table 7 Training parameters for XGBoost regressor

Parameter	Value
objective	reg:squarederror
base score	0.5
booster	gbtree
colsample bylevel	1
colsample bynode	1
colsample bytree	1
gamma ¹	0
learning rate	0.300000012
max delta step	0
max depth	6
min child weight	1
estimators	200
n jobs	12
num parallel tree	1
predictor	auto
random state	0
reg alpha	0
reg lambda	1
scale pos weight	1
subsample	2
tree method	auto

¹Minimum loss reduction required to make a further partition on a leaf node of the tree.

Table 8 Training parameters for XGBoost regressor

Parameter	Value
Maximum iterations	300
Tolerance ¹	1×10^{-3}
alpha 1	1×10^{-6}
alpha 2	1×10^{-6}
lambda 1	1×10^{-6}
lambda 2	1×10^{-6}

¹Tolerance for the stopping criteria.

the model's output as the predicted value of a specific Spotify audio feature. By training the GPT-2 model on a large dataset of Last.fm tags and corresponding Spotify audio features, we aim to learn the complex relationships between these two types of data and to use this knowledge to evaluate the prediction accuracy of Spotify audio features based on Last.fm tags.

To train the GPT-2 model as a regressor, we used a supervised learning approach. Specifically, we collected a large dataset of Last.fm tags and corresponding Spotify audio features, and used this dataset to train the GPT-2 model to predict Spotify audio features from Last.fm tags. We used a mean

squared error loss function to optimize the model's performance during training, and we also employed techniques such as early stopping and learning rate scheduling to prevent overfitting and improve generalization performance. Under development...

4.6 Experiments Execution and Results

Table 9 summaries the experiment results. The table provides RMSE values for each experiment.

4.6.1 Results for Tabular Data Models




Fig. 3 RMSE mean and standard deviation by model and tags/tokens limit.

5 Conclusions

In general, we believe that this novel approach that has the potential to benefit both listeners and researchers. By combining subjective user-generated tags with objective audio features, we can gain new insights into the complex relationship between perception and audio signal in music.

Table 9 Experiment results. Cells values correspond to the RMSE value.

M	Input format	Danceab...	Acoustic...	Energy	Valence	Instrumen...
Base		0.276	0.438	0.329	0.395	0.541
Bayes	100 tags ¹	0.159	0.261	0.197	0.243	0.307
Bayes	1000 tags	0.153	0.253	0.190	0.237	0.299
Bayes	10000 tags	0.152	0.251	0.189	0.236	0.297
Bayes	100 tokens D ²	0.307	0.307	0.238	0.281	0.383
Bayes	10000 tokens D	0.359	0.507	0.394	0.479	0.613
Bayes	1000 tokens D	0.201	0.315	0.249	0.248	0.399
Bayes	100 tokens O ³	0.191	0.305	0.237	0.282	0.376
Bayes	1000 tokens O	0.237	0.343	0.276	0.339	0.428
Bayes	10000 tokens O	0.237	0.343	0.276	0.339	0.428
Bayes	100 tokens TC ⁴	0.191	0.304	0.236	0.281	0.380
Bayes	1000 tokens TC	0.202	0.320	0.247	0.248	0.404
Bayes	10000 tokens TC	0.234	0.341	0.274	0.321	0.430
Tree	100 tags	0.154	0.257	0.188	0.240	0.302
Tree	1000 tags	0.149	0.249	0.184	0.236	0.291
Tree	10000 tags	0.148	0.250	0.181	0.235	0.290
Tree	100 tokens D	0.274	0.274	0.212	0.256	0.330
Tree	1000 tokens D	0.173	0.278	0.215	0.220	0.339
Tree	10000 tokens D	0.172	0.276	0.217	0.266	0.342
Tree	100 tokens O	0.179	0.294	0.225	0.271	0.350
Tree	1000 tokens O	0.182	0.294	0.224	0.270	0.353
Tree	10000 tokens O	0.182	0.294	0.225	0.267	0.353
Tree	100 tokens TC	0.172	0.280	0.211	0.262	0.342
Tree	1000 tokens TC	0.174	0.282	0.215	0.220	0.344
Tree	10000 tokens TC	0.175	0.281	0.214	0.270	0.345
GPT	Duplicated ⁵	0.157	0.244	0.193	0.245	0.322
GPT	Ordered ⁶	0.149	0.237	0.188	0.235	0.297
GPT	Tags,Counts ⁷	0.145	0.237	0.187	0.233	0.301

¹Tags in tabular format. Given $(rock, 3)$, the cell in the *rock* column contains 3.

²Duplicated tokens in tabular format. Tags $(rock, 3)$, $(pop, 2)$, are converted to the "**rock, rock, rock, pop, pop**" string, which a tokenizer converts to the list of input tokens (e.g. [101,1005,16588,1005,2531, ...]). These tokens are passed to the model in tabular format. Columns are $token_1, token_2, \dots, token_N$.

³Ordered tokens in tabular format. Tags $(rock, 3)$, $(pop, 2)$, are converted to the "**rock, pop**" string, which a tokenizer converts to the list of input tokens (e.g. [101,1005,16588,1005,2531, ...]). These tokens are passed to the model in tabular format. Columns are $token_1, token_2, \dots, token_N$.

⁴Tokens in tabular format from tags and counts. Tags $(rock, 3)$, $(pop, 2)$, are converted to the "**'rock' 3, 'pop' 2**" string, which a tokenizer converts to the list of input tokens (e.g. [101,1005,16588,1005,2531, ...]). These tokens are passed to the model in tabular format. Columns are $token_1, token_2, \dots, token_N$.

⁵String. Given tags $(rock, 3)$, $(pop, 2)$, input is formatted as "**rock, rock, rock, pop, pop**".

⁶String. Given tags $(rock, 3)$, $(pop, 2)$, input is formatted as "**rock, pop**".

⁷String. Given tags $(rock, 3)$, $(pop, 2)$, input is formatted as "**'rock' 3, 'pop' 2**".

Our approach also presents limitations. One limitation is the assumption of a strong relationship between Last.fm tags and Spotify features, which may




Fig. 4 RMSE mean and standard deviation by model and audio feature.

not be true in all cases. Future work could explore other sources of input values, possibly related to the user context, to improve the accuracy of the predictions.

TODO

6 Acknowledgments

TODO

References

- [1] Ramirez, J., Flores, M.J.: Machine learning for music genre: multi-faceted review and experimentation with audioset. *Journal of Intelligent Information Systems* **55**(3), 469–499 (2020) <https://doi.org/10.1007/s10844-019-00582-9>
- [2] Laurier, C., Sordo, M., Serra, J., Herrera, P.: Music mood representations from social tags. In: *ISMIR*, pp. 381–386 (2009)
- [3] Çano, E., Morisio, M., *et al.*: Music mood dataset creation based on last. fm tags. In: *International Conference on Artificial Intelligence and Applications*, Vienna, Austria, pp. 15–26 (2017)




Fig. 5 RMSE mean and standard deviation by model and input type (tag probabilities or tokens).

- [4] Bodó, Z., Szilágyi, E.: Connecting the last. fm dataset to lyricwiki and musicbrainz. lyrics-based experiments in genre classification. *Acta Universitatis Sapientiae, Informatica* **10**(2), 158–182 (2018)
- [5] Bertin-Mahieux, T., Ellis, D.P.W., Whitman, B., Lamere, P.: The million song dataset. In: *Proceedings of the 12th International Conference on Music Information Retrieval (ISMIR)* (2011)
- [6] Wang, Y., Horvát, E.-Á.: Gender differences in the global music industry: Evidence from musicbrainz and the echo nest. In: *Proceedings of the International AAAI Conference on Web and Social Media*, vol. 13, pp. 517–526 (2019)
- [7] Jamdar, A., Abraham, J., Khanna, K., Dubey, R.: Emotion analysis of songs based on lyrical and audio features. *arXiv preprint arXiv:1506.05012* (2015)
- [8] Benzi, K., Kalofolias, V., Bresson, X., Vandergheynst, P.: Song recommendation with non-negative matrix factorization and graph total variation. In: *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 2439–2443 (2016). Ieee




Fig. 6 RMSE mean and standard deviation by audio feature and input type (tag probabilities or tokens).

- [9] Panda, R., Redinho, H., Gonçalves, C., Malheiro, R., Paiva, R.P.: How does the spotify api compare to the music emotion recognition state-of-the-art? In: 18th Sound and Music Computing Conference (SMC 2021), pp. 238–245 (2021)
- [10] Pinter, A.T., Paul, J.M., Smith, J., Brubaker, J.R.: P4kxspotify: A dataset of pitchfork music reviews and spotify musical features. In: Proceedings of the International AAAI Conference on Web and Social Media, vol. 14, pp. 895–902 (2020)
- [11] Ramirez-Castillo, J., Flores, M.J., Nicholson, A.E.: User-centric music recommendations. In: 16th Bayesian Modelling Applications Workshop, Conference on Uncertainty in Artificial Intelligence, Eindhoven, The Netherlands) (2022)
- [12] xgboost: Xgboost. <https://xgboost.readthedocs.io/en/stable/tutorials/model.html>
- [13] Tipping, M.E.: Sparse bayesian learning and the relevance vector machine. J. Mach. Learn. Res. **1**, 211–244 (2001)